

## Task 3: Machine translation using Transformer model

This experiment implemented a Transformer-based encoder-decoder model for English-to-Tamil machine translation. The goal was to explore whether attention-based architectures—especially Transformers—could yield improved translation quality and faster training compared to previous RNN-based models, while maintaining the same input and preprocessing pipeline. The experiment reused previously preprocessed sentence pairs, GloVe embeddings for English, and IndicBERT tokenization for Tamil.

### Experiment:

- **Source: English**
- Tokenization: A custom Vocab class was used to tokenize English sentences into word-level tokens.
- Embeddings: 100-dimensional GloVe embeddings were used.
- Each word token was mapped to a 100-dimensional GloVe vector.
- Pre-trained GloVe embeddings were loaded and frozen during training.
- A special <pad> token was assigned a zero vector.
- Unknown words were mapped to a special <unk> embedding.
- Embedding matrix size: (vocab\_size, 100), where vocab\_size  $\approx$  number of unique English tokens.
- Target: Tamil
- Target Tokenization: IndicBERT tokenizer was used.
- Operates at the subword level using SentencePiece-like tokenization.
- Token IDs were generated directly from the tokenizer.
- IndicBERT's tokenizer vocabulary size was large:  $\approx$  200,000 tokens.
- Embeddings: Tokens were converted to embeddings through a learnable embedding layer of the Transformer decoder (not pre-trained).
- Each token was represented by a learnable 100-dimensional vector.
- Positional embeddings were added to this embedding.
- Output token IDs were predicted using a linear projection layer followed by a log-softmax across the entire 200k vocabulary.
- Model Architecture: Transformer (Encoder-Decoder)
- The model followed the standard Transformer architecture introduced in Vaswani et al. (2017), with a few modifications for memory and performance efficiency.

- Encoder (for English Input):
- Number of layers: 2
- Self-Attention: Multi-head self-attention applied to GloVe embeddings with added positional encodings.
- Embedding dimension (d\_model): 100
- Number of attention heads (nhead): 2
- Each head operates on 50-dimensional projections (d\_model / nhead)
- Feedforward network dimension (dim\_feedforward): 128
- Positional Encoding: Learnable positional embeddings of shape (max\_len, d\_model)
- Dropout: 0.1 applied in attention and FF layers
- LayerNorm: Used after attention and feedforward layers
- Decoder (for Tamil Output):
- Number of layers: 2
- Input: Tamil token IDs, shifted right for teacher forcing during training
- Token Embedding Layer: Learnable embeddings of shape (200,000, 100)
- Positional Embedding: Learnable positional embeddings of shape (max\_len, d\_model)
- Masked Multi-head Self-Attention: Prevents attending to future positions
- Cross-Attention: Allows the decoder to attend to encoder output (i.e., context from the English sentence)
- Feedforward network: Same as in encoder (128 hidden units)
- Dropout: 0.1
- LayerNorm: Applied similarly
- Output Layer:
- A linear layer maps decoder outputs of shape (batch\_size, tgt\_seq\_len, d\_model) to logits of shape (batch\_size, tgt\_seq\_len, vocab\_size), where vocab\_size = 200,000.
- Logits are passed through log\_softmax (via nn.NLLLoss during training).
- Loss function: nn.NLLLoss(ignore\_index=pad\_token\_id, label\_smoothing=0.1)

### ***Examples from training set***

EN:The financial assistance is provided to newly married couples of whom one spouse should be from Scheduled caste, Scheduled tribe and the other from a different Community.

T:பிரிவு-1 புதுமணத் தம்பதியரில் ஒருவர் ஆதிதிராவிடர் அல்லது பழங்குடியினராக இருந்து பிற இனத்தவரை மணந்து கொண்டால் நிதியுதவி வழங்கப்படும்.

M:இந்தப்பட்டரகளக்க மன்னரகளக்கப்பட்டவரகளக்கப்பட்டவரகளக்க வண்டம்.

EN:Congress government

HI-T:காங்கிரஸ் அரசின் பலம்

HI-M:அரசயல கங்கரஸ்

EN:Then there was a loud thud.

HI-T:அப்போது கரடி ஒன்று உறுமும் சத்தம் கேட்டது.

HI-M:பன்னர அதல அதலம அதகரத்த.

EN:Vacancy: Manager

HI-T:பணியிடம்: சண்டகர்

HI-M:வண்டம்

- The Transformer completed 5 full epochs in just 25 minutes, a massive leap in training speed compared to the RNN variant, which took 2.5 hours per epoch, totaling 21 hours across 7 runs. This drastic improvement is due to:
- Parallel computation in self-attention layers (vs. sequential processing in RNNs).
  - GPU-optimized matrix operations and feedforward layers in Transformer blocks.
  - Faster gradient flow and efficient backpropagation.
- The ability to leverage hardware acceleration made Transformer training highly feasible and practical.
- Despite having similar training settings, the Transformer achieved:
- Final Test Loss: 3.0348
  - Test Perplexity: 20.80

BLEU Scores:

B@1: 0.2507

B@2: 0.1469

B@3: 0.0924

B@4: 0.0564

These BLEU scores, although modest in absolute terms, are significantly better than the RNN model, which barely crossed 0.148 even after extensive training.

- Crucially, the output translations were far more structured, meaningful, and close to the true targets – a sharp contrast to the repetitive and broken outputs of the RNN model.
- The model output, though not perfect, shows a recognizable structure, some correctly translated fragments, and a far better grasp of syntax than the RNN version which typically just repeated [CLS] or nonsensical patterns.
- EN: There are 15 countries.
- HI-T: இதில் 15 நாடுகள் இடம் பெற்றுள்ளன.
- HI-M: இதல 20 நடகளல இரககறத.
- ➤ Again, the Transformer predicts something semantically related ("countries" is somewhat reflected), with errors in number and spelling – yet it's a reasonable approximation.

## Conclusion

The Transformer model not only trained exponentially faster but also produced substantially better translations, both quantitatively (loss, BLEU) and qualitatively (output structure). It proved to be a more viable and scalable choice for neural machine translation, even under the constraints of subword token noise and large vocabulary sizes – where RNNs completely failed to learn.